

Automatic Differentiation

In this talk:

- ▶ How to calculate derivatives inside computer?
- ▶ Forward and reverse mode automatic differentiation
- ▶ Computational graphs and evaluation traces
- ▶ Engineering trade-offs

Annada Behera


June 10, 2022

Computing derivatives: Numerical approximation

For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the gradient can be computed by Newton's method,

$$\nabla f(x) = \lim_{h \rightarrow 0} \frac{f(x + e_i h) - f(x)}{h} \quad (1)$$

where e_i is the unit vector in i -th dimension.




Computing derivatives: Numerical approximation

For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the gradient can be computed by Newton's method,


$$\nabla f(x) = \lim_{h \rightarrow 0} \frac{f(x + e_i h) - f(x)}{h} \quad (1)$$

where e_i is the unit vector in i -th dimension.



Idea: Use arbitrary-precision arithmetic. Takes $O(nk^2)$ time for n variables with k digits of precision. **And this is still just an approximation!**

Symbolic differentiation



$$y_1 = \sigma(w_{00}x_0) + \sigma(w_{10}x_1) + \sigma(w_{20}x_2) + \sum_j \sigma\left(\sum_i w_{ij}x_i\right) \quad (2)$$

where σ is any non-linear function. Or, if $x = x_i$, $y = y_i$, $W^0 = w_{ij}$, and $W^1 = w_{ij}$, then,

$$y = W^1 \sigma(W^0 x) \quad (3)$$

Consider the product rule,

$$\frac{\partial fg}{\partial x} = \frac{\partial f}{\partial x} \cdot g + f \cdot \frac{\partial g}{\partial x} \quad (4)$$


This method of computing derivative of n variables takes $O(2^n)$ time!

Computational graph and evaluation traces


Consider the equation,

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)][x_1/x_2 - \exp(x_2)] \quad (5)$$

The computational graph (DAG):




The decomposition of calculations into elementary steps forms the *evaluation trace*. Say, $x = 1.5$ and $x_2 = 0.5$, then the evaluation trace will look like,



These intermediate values are called the *primal trace*.

Forward mode: Accumulating the *tangent trace*



For $x_1 = 1.5$ and $x_2 = 0.5$, the primal and tangent trace with respect to x_1 will be,

$$u_{-1} = x_1$$

$$u_0 = x_2$$

$$u_1 = u_{-1}/u_0$$

$$u_2 = \sin(u_1)$$

$$u_3 = \exp(u_0)$$

$$u_4 = u_1 - u_3$$

$$u_5 = u_2 + u_4$$

$$u_6 = u_5 u_4$$

$$\dot{u}_{-1} = \frac{\partial u_{-1}}{\partial x_1} = \partial_{x_1} x_1 = 1$$


$$\dot{u}_0 = \partial_{x_1} x_2 = 0$$

$$\begin{aligned} \dot{u}_1 &= \partial(u_{-1}/u_0) = u_{-1} \partial_{x_1} (1/u_0) + \partial_{x_1} u_{-1} \cdot \frac{1}{u_0} \\ &= 0 + 1/0.5 = 2 \end{aligned}$$

$$\dot{u}_2 = \partial_{x_1} \sin(u_1) = \cos(u_1) \dot{u}_1 = \cos(3.0) \times 2 = -1.98$$

$$\dot{u}_3 = \partial_{x_1} \exp(u_0) = \exp(u_0) \dot{u}_0 = 0$$

Accumulating tangent trace (contd.)



Mathematica:

```
In[1]:= D[(Sin[x/0.5]+x/0.5-Exp[0.5])*(x/0.5-Exp[0.5]),x]/.x->1.5
```

```
Out[0]= 3.01184
```

We can do the same and get the derivative with respect to x_2 . Therefore for any function, $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, it takes $O(n)$ time.

Forward mode is a great choice for $f : \mathbb{R} \rightarrow \mathbb{R}^m$ since, takes **constant** time but it is a poor choice for $f : \mathbb{R}^n \rightarrow \mathbb{R}$.

How to code?

A *dual number* takes the form,

$$x = a + b\epsilon \quad \text{where, } \epsilon^2 = 0 \text{ and } a, b \in \mathbb{R} \quad (6)$$

The component-wise addition and multiplication is given by,

$$(a + b\epsilon)(c + d\epsilon) = ac + (ad + bc)\epsilon \quad (7)$$

It is not hard to see that dual numbers form a *commutative algebra* over two dimension.

The Taylor series expansion of any arbitrary function f at the dual number $a + \epsilon$ is,

$$f(a + b\epsilon) = \sum_{n=0}^{\infty} \frac{f^{(n)}(a)b^n\epsilon^n}{n!} = f(a) + bf'(a)\epsilon \quad (8)$$

Evaluate any function at $a + \epsilon$, the co-efficient of ϵ gives the derivative of the function.

Forward mode autodiff: an example

- Product rule:

$$f(a + \epsilon)g(a + \epsilon) = [f(a) + \dot{f}(a)\epsilon][g(a) + \dot{g}(a)\epsilon] \quad (9)$$

$$= f(a)g(a) + [\dot{f}(a)g(a) + f(a)\dot{g}(a)]\epsilon \quad (10)$$

- Chain rule:

$$f(g(a + \epsilon)) = f(g(a) + \dot{g}(a)\epsilon) \quad (11)$$

$$= f(g(a)) + \dot{f}(g(a))\dot{g}(a)\epsilon \quad (12)$$

- An example: What is the derivative of $f(x) = kx + \sin(x)$?

$$f(a + \epsilon) = k \cdot (a + \epsilon) + \sin(a + \epsilon) \quad (13)$$

$$= ka + k\epsilon + \sin(a) \cos(\epsilon) + \cos(a) \sin(\epsilon) \quad (14)$$

$$= \underbrace{ka + \sin(a)}_{\text{primal trace}} + \underbrace{\epsilon(k + \cos(a))}_{\text{tangent trace}} \quad (15)$$

Reverse mode autodiff: Propagate from the end


Chain rule: For any arbitrary function, $y = f(x(t))$,

$$\partial_t y = \partial_x y \cdot \partial_t x \quad (16)$$

Consider the same function again,

$$y = [\sin(x_1/x_2) + x_1/x_2 - \exp(x_2)][x_1/x_2 - \exp(x_2)] \quad (17)$$


along with it's variables and DA computational graph,



Define the *adjoint* \bar{u}_i as,

$$\bar{u}_i = \partial_{u_i} y \quad (18)$$

Reverse mode autodiff (contd.)



$$u_6 = u_5 u_4 = 2.016$$

$$u_5 = u_2 + u_4 = 1.492$$

$$u_4 = u_1 - u_3 = 1.351$$

$$u_3 = \exp(u_0) = 1.648$$

$$u_2 = \sin(u_1) = 0.141$$

$$u_1 = u_{-1}/u_0 = 3$$

$$u_0 = x_2 = 0.5$$


$$u_{-1} = x_1 = 1.5$$

$$\partial_{u_6} y = 1$$

$$\partial_{u_5} y = \partial_{u_6} y \cdot \partial_{u_5} u_6 = 1.351$$

$$\begin{aligned} \partial_{u_4} y &= \partial_{u_6} y \cdot \partial_{u_4} u_6 + \partial_{u_5} y \cdot \partial_{u_4} u_5 \\ &= 1.351 + 1.492 = 2.844 \end{aligned}$$

Reverse mode autodiff (contd.)



$$\bar{x}_1 = \partial_{x_1} y = 3.012 \quad \text{and,} \quad \bar{x}_2 = \partial_{x_2} y = -13.724 \quad (19)$$

- ▶ For a function, $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, the reverse mode autodiff **takes $O(m)$ time!**
- ▶ Great choice for $f : \mathbb{R}^n \rightarrow \mathbb{R}$ but poor choice for $g : \mathbb{R} \rightarrow \mathbb{R}^m$.
- ▶ In neural networks, the loss function is, $L : \mathbb{R}^n \rightarrow \mathbb{R}$. Therefore the gradients can be computed in **constant** time.
- ▶ We are essentially trading space for time.

Engineering trade-offs: Static graphs

Import the Tensorflow 1.0 package:

```
import tensorflow as tf
```


Creating variables:

```
k = tf.Variable(2.0, name='k')  
b = tf.Variable(2.0, name='b')  
c = tf.Variable(1.0, name='c')
```


Doing calculations:

```
d = tf.add(b, c, name='d')  
e = tf.add(c, k, name='e')  
a = tf.multiply(d, e, name='a')
```

Computational graph:



Dynamic graphs



```
from torch import nn
import torch
class calc(nn.Module):
    def __init__(self):
        super().__init__()
        self.k = torch.tensor([2],
                               requires_grad=True)
    def forward(self, b, c):
        d = b+c
        e = c+self.k
        return e*d

b = 2
c = 1
a = calc(b, c)
```

- ▶ All computations must be a part of nn.Module.
- ▶ All parameters must be top level variables in the constructor.
- ▶ All computation is stored on a 'tape' during the forward pass.
- ▶ Hence, dynamic graphs are slow.

Summary

- ▶ Numerical differentiation is slow and inaccurate.
- ▶ Symbolic differentiation is extremely slow and hard to code.
- ▶ Algorithmic/automatic differentiation is fast and exact.
- ▶ Forward-mode autodiff is better for $f : \mathbb{R} \rightarrow \mathbb{R}^m$.
- ▶ Reverse-mode autodiff is better for $f : \mathbb{R}^m \rightarrow \mathbb{R}$.
- ▶ Forward mode takes less space but needs more time.
- ▶ Reverse mode takes more space and needs less time.
- ▶ Neural network libraries use reverse mode autodiff.
- ▶ Dynamic computation graph is slow but easy to write, e.g, PyTorch.
- ▶ Static computation graph is fast but hard to write, e.g, Tensorflow.
- ▶ JAX and Julia make trade-offs in between static and dynamic graphs using *jit*.

Thank you.